
Introduction


In this notebook, we implement the tiny version of [YOLOv4](#) for training on your own dataset, [YOLOv4 tiny](#).

We also recommend reading our blog post on [Training YOLOv4 on custom data](#) side by side.

We will take the following steps to implement YOLOv4 on our custom data:

- Configure our GPU environment on Google Colab
- Install the Darknet YOLOv4 training environment
- Download our custom dataset for YOLOv4 and set up directories
- Configure a custom YOLOv4 training config file for Darknet
- Train our custom YOLOv4 object detector
- Reload YOLOv4 trained weights and make inference on test images

When you are done you will have a custom detector that you can use. It will make inference like this:



✓ Configuring cuDNN on Colab for YOLOv4

```
# CUDA: Let's check that Nvidia CUDA drivers are already pre-installed and which
!/usr/local/cuda/bin/nvcc --version
# We need to install the correct cuDNN according to this output
```

```
⇒ nvcc: NVIDIA (R) Cuda compiler driver
   Copyright (c) 2005-2024 NVIDIA Corporation
   Built on Thu_Jun__6_02:18:23_PDT_2024
   Cuda compilation tools, release 12.5, V12.5.82
   Build cuda_12.5.r12.5/compiler.34385749_0
```

```
#take a look at the kind of GPU we have
!nvidia-smi
```

```
➡ Sat Mar 8 01:20:17 2025
```

NVIDIA-SMI 550.54.15			Driver Version: 550.54.15			CUDA Version: 12.5.2		
GPU	Name		Persistence-M	Bus-Id	Disp.A	Volatile		
Fan	Temp	Perf	Pwr:Usage/Cap		Memory-Usage	GPU-Util		
0	Tesla T4		Off	00000000:00:04.0	Off			
N/A	56C	P8	12W / 70W		0MiB / 15360MiB	0%		
Processes:								
GPU	GI	CI	PID	Type	Process name			
	ID	ID						
No running processes found								

```
# This cell ensures you have the correct architecture for your respective GPU
# If you command is not found, look through these GPUs, find the respective
# GPU and add them to the archTypes dictionary

# Tesla V100
# ARCH= -gencode arch=compute_70,code=[sm_70,compute_70]

# Tesla K80
# ARCH= -gencode arch=compute_37,code=sm_37

# GeForce RTX 2080 Ti, RTX 2080, RTX 2070, Quadro RTX 8000, Quadro RTX 6000, Quad
# ARCH= -gencode arch=compute_75,code=[sm_75,compute_75]

# Jetson XAVIER
# ARCH= -gencode arch=compute_72,code=[sm_72,compute_72]

# GTX 1080, GTX 1070, GTX 1060, GTX 1050, GTX 1030, Titan Xp, Tesla P40, Tesla P4
# ARCH= -gencode arch=compute_61,code=sm_61

# GP100/Tesla P100 - DGX-1
# ARCH= -gencode arch=compute_60,code=sm_60

# For Jetson TX1, Tegra X1, DRIVE CX, DRIVE PX - uncomment:
# ARCH= -gencode arch=compute_53,code=[sm_53,compute_53]
```

```

# For Jetson Tx2 or Drive-PX2 uncomment:
# ARCH= -gencode arch=compute_62,code=[sm_62,compute_62]
import os
os.environ['GPU_TYPE'] = str(os.popen('nvidia-smi --query-gpu=name --format=csv,nl

def getGPUArch(argument):
    try:
        argument = argument.strip()
        # All Colab GPUs
        archTypes = {
            "Tesla V100-SXM2-16GB": "-gencode arch=compute_70,code=[sm_70,compute_70]",
            "Tesla K80": "-gencode arch=compute_37,code=sm_37",
            "Tesla T4": "-gencode arch=compute_75,code=[sm_75,compute_75]",
            "Tesla P40": "-gencode arch=compute_61,code=sm_61",
            "Tesla P4": "-gencode arch=compute_61,code=sm_61",
            "Tesla P100-PCI-E-16GB": "-gencode arch=compute_60,code=sm_60"

        }
        return archTypes[argument]
    except KeyError:
        return "GPU must be added to GPU Commands"
os.environ['ARCH_VALUE'] = getGPUArch(os.environ['GPU_TYPE'])

print("GPU Type: " + os.environ['GPU_TYPE'])
print("ARCH Value: " + os.environ['ARCH_VALUE'])

```

GPU Type: Tesla T4

ARCH Value: -gencode arch=compute_75,code=[sm_75,compute_75]

✓ Installing Darknet for YOLOv4 on Colab

```

%cd /content/
%rm -rf darknet

```

/content

```
#we clone the fork of darknet maintained by roboflow
#small changes have been made to configure darknet for training
!git clone https://github.com/roboflow-ai/darknet.git
```

```
➦ Cloning into 'darknet'...
remote: Enumerating objects: 13289, done.
remote: Total 13289 (delta 0), reused 0 (delta 0), pack-reused 13289 (from 1)
Receiving objects: 100% (13289/13289), 12.17 MiB | 4.40 MiB/s, done.
^C
```

```
#install environment from the Makefile
%cd /content/darknet/
# compute_37, sm_37 for Tesla K80
# compute_75, sm_75 for Tesla T4
# !sed -i 's/ARCH= -gencode arch=compute_60,code=sm_60/ARCH= -gencode arch=compute_75,code=sm_75/g' Makefile
```

```
#install environment from the Makefile
#note if you are on Colab Pro this works on a P100 GPU
#if you are on Colab free, you may need to change the Makefile for the K80 GPU
#this goes for any GPU, you need to change the Makefile to inform darknet which GPU to use
!sed -i 's/OPENCV=0/OPENCV=1/g' Makefile
!sed -i 's/GPU=0/GPU=1/g' Makefile
!sed -i 's/CUDNN=0/CUDNN=1/g' Makefile
!sed -i "s/ARCH= -gencode arch=compute_60,code=sm_60/ARCH= ${ARCH_VALUE}/g" Makefile
!make
```

```
➦ [Errno 2] No such file or directory: '/content/darknet/'
/content
sed: can't read Makefile: No such file or directory
sed: can't read Makefile: No such file or directory
sed: can't read Makefile: No such file or directory
sed: can't read Makefile: No such file or directory
make: *** No targets specified and no makefile found. Stop.
```

```

#download the newly released yolov4-tiny weights
%cd /content/darknet
!wget https://github.com/AlexeyAB/darknet/releases/download/darknet_yolo_v4_pre/yolov4-tiny.weights
!wget https://github.com/AlexeyAB/darknet/releases/download/darknet_yolo_v4_pre/yolov4-tiny.conv.29

[Errno 2] No such file or directory: '/content/darknet'
/content
--2025-03-08 01:20:23-- https://github.com/AlexeyAB/darknet/releases/download/darknet_yolo_v4_pre/yolov4-tiny.weights
Resolving github.com (github.com)... 140.82.113.3
Connecting to github.com (github.com)|140.82.113.3|:443... connected.
HTTP request sent, awaiting response... 302 Found
Location: https://objects.githubusercontent.com/github-production-release-asset/29613631/29613631
--2025-03-08 01:20:23-- https://objects.githubusercontent.com/github-production-release-asset/29613631/29613631
Resolving objects.githubusercontent.com (objects.githubusercontent.com)... 185.199.108.133
Connecting to objects.githubusercontent.com (objects.githubusercontent.com)|185.199.108.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 24251276 (23M) [application/octet-stream]
Saving to: 'yolov4-tiny.weights'

yolov4-tiny.weights 100%[=====>] 23.13M 64.8MB/s in 0.4s

2025-03-08 01:20:24 (64.8 MB/s) - 'yolov4-tiny.weights' saved [24251276/24251276]

--2025-03-08 01:20:24-- https://github.com/AlexeyAB/darknet/releases/download/darknet_yolo_v4_pre/yolov4-tiny.conv.29
Resolving github.com (github.com)... 140.82.112.4
Connecting to github.com (github.com)|140.82.112.4|:443... connected.
HTTP request sent, awaiting response... 302 Found
Location: https://objects.githubusercontent.com/github-production-release-asset/29613631/29613631
--2025-03-08 01:20:24-- https://objects.githubusercontent.com/github-production-release-asset/29613631/29613631
Resolving objects.githubusercontent.com (objects.githubusercontent.com)... 185.199.108.133
Connecting to objects.githubusercontent.com (objects.githubusercontent.com)|185.199.108.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 19789716 (19M) [application/octet-stream]
Saving to: 'yolov4-tiny.conv.29'

yolov4-tiny.conv.29 100%[=====>] 18.87M 80.5MB/s in 0.2s

2025-03-08 01:20:25 (80.5 MB/s) - 'yolov4-tiny.conv.29' saved [19789716/19789716]

```

✓ Set up Custom Dataset for YOLOv4

We'll use Roboflow to convert our dataset from any format to the YOLO Darknet format.

1. To do so, create a free [Roboflow account](#).
2. Upload your images and their annotations (in any format: VOC XML, COCO JSON, TensorFlow CSV, etc).
3. Apply preprocessing and augmentation steps you may like. We recommend at least auto-orient and a resize to 416x416. Generate your dataset.
4. Export your dataset in the **YOLO Darknet format**.
5. Copy your download link, and paste it below.

See our [blog post](#) for greater detail.

In this example, I used the open source [BCCD Dataset](#). (You can fork it to your Roboflow account to follow along.)

```
#follow the link below to get your download code from from Roboflow
!pip install -q roboflow
from roboflow import Roboflow
rf = Roboflow(model_format="darknet", notebook="roboflow-yolov4-tiny")
```

```
83.1/83.1 kB 2.6 MB/s eta 0:00:00
66.8/66.8 kB 6.3 MB/s eta 0:00:00
49.9/49.9 MB 19.6 MB/s eta 0:00:00
```

RuntimeError Traceback (most recent call last)

```
<ipython-input-8-05fac0eb4601> in <cell line: 0>()
      2 get_ipython().system('pip install -q roboflow')
      3 from roboflow import Roboflow
----> 4 rf = Roboflow(model_format="darknet", notebook="roboflow-yolov4-
      tiny")
```

2 frames

```
/usr/local/lib/python3.11/dist-packages/roboflow/__init__.py in
check_key(api_key, model, notebook, num_retries)
      21 def check_key(api_key, model, notebook, num_retries=0):
      22     if not isinstance(api_key, str):
----> 23         raise RuntimeError(
      24             "API Key is of Incorrect Type \n Expected Type: " +
str(str) + "\n Input Type: " + str(type(api_key))
      25         )
```

```
RuntimeError: API Key is of Incorrect Type
Expected Type: <class 'str'>
Input Type: <class 'NoneType'>
```

Next steps: [Explain error](#)

```
# from roboflow import Roboflow
# rf = Roboflow(api_key="YOUR_API_KEY")
# project = rf.workspace().project("YOUR_PROJECT")
# dataset = project.version("YOUR_VERSION").download("darknet")
```



```

#Set up training file directories for custom dataset
%cd /content/darknet/
%cp {dataset.location}/train/_darknet.labels data/obj.names
%mkdir data/obj
#copy image and labels
%cp {dataset.location}/train/*.jpg data/obj/
%cp {dataset.location}/valid/*.jpg data/obj/

%cp {dataset.location}/train/*.txt data/obj/
%cp {dataset.location}/valid/*.txt data/obj/

with open('data/obj.data', 'w') as out:
    out.write('classes = 3\n')
    out.write('train = data/train.txt\n')
    out.write('valid = data/valid.txt\n')
    out.write('names = data/obj.names\n')
    out.write('backup = backup/')

#write train file (just the image list)
import os

with open('data/train.txt', 'w') as out:
    for img in [f for f in os.listdir(dataset.location + '/train') if f.endswith('.jpg')]:
        out.write('data/obj/' + img + '\n')

#write the valid file (just the image list)
import os

with open('data/valid.txt', 'w') as out:
    for img in [f for f in os.listdir(dataset.location + '/valid') if f.endswith('.jpg')]:
        out.write('data/obj/' + img + '\n')

```

✓ Write Custom Training Config for YOLOv4

```

#we build config dynamically based on number of classes
#we build iteratively from base config files. This is the same file shape as cfg/
def file_len(fname):
    with open(fname) as f:
        for i, l in enumerate(f):
            pass
    return i + 1

num_classes = file_len(dataset.location + '/train/_darknet.labels')
max_batches = num_classes*2000
steps1 = .8 * max_batches
steps2 = .9 * max_batches
steps_str = str(steps1)+' '+str(steps2)
num_filters = (num_classes + 5) * 3

print("writing config for a custom YOLOv4 detector detecting number of classes: ")

#Instructions from the darknet repo
#change line max_batches to (classes*2000 but not less than number of training images)
#change line steps to 80% and 90% of max_batches, f.e. steps=4800,5400
if os.path.exists('./cfg/custom-yolov4-tiny-detector.cfg'): os.remove('./cfg/custom-yolov4-tiny-detector.cfg')

#customize iPython writefile so we can write variables
from IPython.core.magic import register_line_cell_magic

@register_line_cell_magic
def writetemplate(line, cell):
    with open(line, 'w') as f:
        f.write(cell.format(**globals()))

%writetemplate ./cfg/custom-yolov4-tiny-detector.cfg
[net]
# Testing
#batch=1
#subdivisions=1
# Training
batch=64
subdivisions=24
width=416
height=416
channels=3
momentum=0.9
decay=0.0005
angle=0
saturation = 1.5

```

```
exposure = 1.5
hue=.1
```

```
learning_rate=0.00261
burn_in=1000
max_batches = {max_batches}
policy=steps
steps={steps_str}
scales=.1,.1
```

```
[convolutional]
batch_normalize=1
filters=32
size=3
stride=2
pad=1
activation=leaky
```

```
[convolutional]
batch_normalize=1
filters=64
size=3
stride=2
pad=1
activation=leaky
```

```
[convolutional]
batch_normalize=1
filters=64
size=3
stride=1
pad=1
activation=leaky
```

```
[route]
layers=-1
groups=2
group_id=1
```

```
[convolutional]
batch_normalize=1
filters=32
size=3
stride=1
pad=1
activation=leaky
```

```
[convolutional]
```

```
batch_normalize=1
filters=32
size=3
stride=1
pad=1
activation=leaky
```

```
[route]
layers = -1,-2
```

```
[convolutional]
batch_normalize=1
filters=64
size=1
stride=1
pad=1
activation=leaky
```

```
[route]
layers = -6,-1
```

```
[maxpool]
size=2
stride=2
```

```
[convolutional]
batch_normalize=1
filters=128
size=3
stride=1
pad=1
activation=leaky
```

```
[route]
layers=-1
groups=2
group_id=1
```

```
[convolutional]
batch_normalize=1
filters=64
size=3
stride=1
pad=1
activation=leaky
```

```
[convolutional]
batch_normalize=1
filters=64
```

```
filters=64  
size=3  
stride=1  
pad=1  
activation=leaky
```

```
[route]  
layers = -1,-2
```

```
[convolutional]  
batch_normalize=1  
filters=128  
size=1  
stride=1  
pad=1  
activation=leaky
```

```
[route]  
layers = -6,-1
```

```
[maxpool]  
size=2  
stride=2
```

```
[convolutional]  
batch_normalize=1  
filters=256  
size=3  
stride=1  
pad=1  
activation=leaky
```

```
[route]  
layers=-1  
groups=2  
group_id=1
```

```
[convolutional]  
batch_normalize=1  
filters=128  
size=3  
stride=1  
pad=1  
activation=leaky
```

```
[convolutional]  
batch_normalize=1  
filters=128  
size=3
```

```
stride=1
pad=1
activation=leaky
```

```
[route]
layers = -1,-2
```

```
[convolutional]
batch_normalize=1
filters=256
size=1
stride=1
pad=1
activation=leaky
```

```
[route]
layers = -6,-1
```

```
[maxpool]
size=2
stride=2
```

```
[convolutional]
batch_normalize=1
filters=512
size=3
stride=1
pad=1
activation=leaky
```

```
#####
```

```
[convolutional]
batch_normalize=1
filters=256
size=1
stride=1
pad=1
activation=leaky
```

```
[convolutional]
batch_normalize=1
filters=512
size=3
stride=1
pad=1
activation=leaky
```

```
[convolutional]
size=1
stride=1
pad=1
filters={num_filters}
activation=linear
```

```
[yolo]
mask = 3,4,5
anchors = 10,14, 23,27, 37,58, 81,82, 135,169, 344,319
classes={num_classes}
num=6
jitter=.3
scale_x_y = 1.05
cls_normalizer=1.0
iou_normalizer=0.07
iou_loss=ciou
ignore_thresh = .7
truth_thresh = 1
random=0
nms_kind=greedynms
beta_nms=0.6
```

```
[route]
layers = -4
```

```
[convolutional]
batch_normalize=1
filters=128
size=1
stride=1
pad=1
activation=leaky
```

```
[upsample]
stride=2
```

```
[route]
layers = -1, 23
```

```
[convolutional]
batch_normalize=1
filters=256
size=3
stride=1
pad=1
activation=leaky
```

```
activation=leaky
```

```
[convolutional]  
size=1  
stride=1  
pad=1  
filters={num_filters}  
activation=linear
```

```
[yolo]  
mask = 1,2,3  
anchors = 10,14, 23,27, 37,58, 81,82, 135,169, 344,319  
classes={num_classes}  
num=6  
jitter=.3  
scale_x_y = 1.05  
cls_normalizer=1.0  
iou_normalizer=0.07  
iou_loss=ciou  
ignore_thresh = .7  
truth_thresh = 1  
random=0  
nms_kind=greedynms  
beta_nms=0.6
```

```
#here is the file that was just written.  
#you may consider adjusting certain things
```

```
#like the number of subdivisions 64 runs faster but Colab GPU may not be big enough  
#if Colab GPU memory is too small, you will need to adjust subdivisions to 16  
%cat cfg/custom-yolov4-tiny-detector.cfg
```

✓ Train Custom YOLOv4 Detector

```
!./darknet detector train data/obj.data cfg/custom-yolov4-tiny-detector.cfg yolov4-tiny.weights  
#If you get CUDA out of memory adjust subdivisions above!  
#adjust max batches down for shorter training above
```

✓ Infer Custom Objects with Saved YOLOv4 Weights


```

#define utility function
def imShow(path):
    import cv2
    import matplotlib.pyplot as plt
    %matplotlib inline

    image = cv2.imread(path)
    height, width = image.shape[:2]
    resized_image = cv2.resize(image, (3*width, 3*height), interpolation = cv2.INTER_

    fig = plt.gcf()
    fig.set_size_inches(18, 10)
    plt.axis("off")
    #plt.rcParams['figure.figsize'] = [10, 5]
    plt.imshow(cv2.cvtColor(resized_image, cv2.COLOR_BGR2RGB))
    plt.show()

#check if weights have saved yet
#backup houses the last weights for our detector
#(file yolo-obj_last.weights will be saved to the build\darknet\x64\backup\ for e
#(file yolo-obj_xxxx.weights will be saved to the build\darknet\x64\backup\ for e
#After training is complete – get result yolo-obj_final.weights from path build\d
!ls backup
#if it is empty you haven't trained for long enough yet, you need to train for at

#coco.names is hardcoded somewhere in the detector
%cp data/obj.names data/coco.names

#/test has images that we can test our detector on
test_images = [f for f in os.listdir('test') if f.endswith('.jpg')]
import random
img_path = "test/" + random.choice(test_images);

#test out our detector!
!./darknet detect cfg/custom-yolov4-tiny-detector.cfg backup/custom-yolov4-tiny-d
imShow('/content/darknet/predictions.jpg')

```

